

IE 7374 ST: Machine Learning in Engineering

HW-1

Rohit Bokade
NUID: 001280767

March 10, 2021

1 Problem 1

$$p(x) = \begin{cases} \frac{x}{10}, & \text{if } x = 1, 2, 3 \\ 0, & \text{otherwise} \end{cases}$$

$$p(y|x) = \left(\frac{x+1}{2x}\right)^y \left(1 - \frac{x+1}{2x}\right)^{1-y}$$

where $y \in \{0, 1\}$

1.

$$\mathbb{E}(X) = \sum_x xp(x) = \sum_{x=1}^4 x \left(\frac{x}{10}\right) = \boxed{3}$$

$$\text{Var}(X) = \mathbb{E}(X^2) - \mathbb{E}(X)^2 = \sum_{x=1}^4 x^2 \left(\frac{x}{10}\right) - 3^2 = 10 - 9 = \boxed{1}$$

2. By using the conditional equality $p(x, y) = p(y|x)p(x)$, we can find

$$p(x, y) = \left(\frac{x+1}{2x}\right)^y \left(1 - \frac{x+1}{2x}\right)^{1-y} \frac{x}{10}$$

and the marginal $p(y)$ can be found by summing the joint distribution over the domain of x

$$p(y) = \sum_x p(x, y) = \sum_{x=1}^4 \left[\left(\frac{x+1}{2x}\right)^y \left(1 - \frac{x+1}{2x}\right)^{1-y} \frac{x}{10} \right]$$

Next, exploiting the conditional equality again, we can find $p(x|y)$

$$p(x|y) = \frac{p(x, y)}{p(y)} = \frac{\left(\frac{x+1}{2x}\right)^y \left(1 - \frac{x+1}{2x}\right)^{1-y} \frac{x}{10}}{\sum_{x=1}^4 \left[\left(\frac{x+1}{2x}\right)^y \left(1 - \frac{x+1}{2x}\right)^{1-y} \frac{x}{10} \right]}$$

3.

$$\begin{aligned}\mathbb{E}[X|Y=1] &= \sum_{x=1}^4 xp(x|y=1) \\ &= \sum_{x=1}^4 x \left(\frac{\left(\frac{x+1}{2x}\right) \frac{x}{10}}{\sum_{x=1}^4 \left[\left(\frac{x+1}{2x}\right) \frac{x}{10}\right]} \right) = \frac{20}{7}\end{aligned}$$

2 Problem 2: Legal reasoning

1. Let `crime_blood_type` represent that the blood type is the one that was found at the scene of crime and `guilty` represent that the defendant is guilty. Probability that anyone has the blood type is $p(\text{crime_blood_type}) = 0.1$ and therefore probability that the defendant has the blood type `crime_blood_type` is 0.1.

However, there is no evidence provided to assess the dependence between `blood_type` and `guilty`. Therefore, it would be safe and unbiased to assume that $p(\text{guilty}|\text{crime_blood_type}) = p(\text{innocent}|\text{crime_blood_type}) = 0.5$ and $p(\text{guilty}) = p(\text{innocent}) = 0.5$.

Thus,

$$p(\text{crime_blood_type}|\text{innocent}) = \frac{p(\text{innocent}|\text{crime_blood_type})p(\text{crime_blood_type})}{p(\text{innocent})} = \frac{0.5 \cdot 0.1}{0.5} = 0.1$$

- The claim made by the prosecutor is $p(\text{crime_blood_type}|\text{innocent}) = 0.1$ and $p(\text{guilty}) = 0.99$. Thus, $p(\text{crime_blood_type}|\text{innocent}) = 0.1$ might hold true if we make unbiased assumptions, which includes $p(\text{guilty}) = 0.5$ and inferring $p(\text{guilty})$ would be incorrect.
2. The only evidence available is that 1 in 8000 people would have the crime blood type. This would **only imply that the defendant has 1 in 8000 chance of having the crime blood type and not 1 in 8000 chance that the defendant is guilty**. The fault of the defender's argument is making an assumption between the dependence between these two random variables.

3 Problem 3: Maximum Likelihood Estimation (MLE)

$$p(x|\theta_0) = \begin{cases} e^{-\sum_{i=1}^n (x_i - \theta_0)}, & x \geq \theta_0 \\ 0, & \text{otherwise} \end{cases}$$

$$\begin{aligned}L(\theta) &= \prod_{i=1}^n p(x|\theta) \quad \forall x \geq \theta = \prod_{i=1}^n e^{-\sum_{i=1}^n (x_i - \theta)} \\ \log L(\theta) &= \sum_{i=1}^n \log e^{-\sum_{i=1}^n (x_i - \theta)} = - \sum_{i=1}^n x_i + n\theta\end{aligned}$$

Hence,

$$\begin{aligned}\arg \max_{\theta} &= \arg \max_{\theta} \{N\theta - \sum_{i=1}^N X_i\} \text{ s.t. } \theta \leq X_i \\ &= \min \{X_i\} = \boxed{\min\{x_i\} \quad \forall i \in \{1 \cdots n\}}\end{aligned}$$

4 Problem 4: Maximum Likelihood Estimation (MLE)

X is a Bernoulli random variable with probability distribution

1.

$$f(x; p) = \begin{cases} p^x(1-p)^{1-x}, & \text{if } x = \{0, 1\} \\ 0, & \text{otherwise} \end{cases}$$

And the likelihood function can be written as

$$\begin{aligned} L(p) &= \prod_{i=1}^n f(x_i; p) \\ \log L(p) &= \log \sum_{i=1}^n f(x_i; p) = \log \sum_{i=1}^n p^{x_i}(1-p)^{1-x_i} = \log p \sum_{i=1}^n x_i \log(1-p) \sum_{i=1}^n (1-x_i) \\ \frac{\partial \ell(p)}{\partial p} &= \frac{\sum_{i=1}^n x_i}{p} - \frac{\sum_{i=1}^n (1-x_i)}{\log(1-p)} \stackrel{\text{set}}{=} 0 \\ \sum_{i=1}^n x_i - p \sum_{i=1}^n x_i &= p \sum_{i=1}^n (1-x_i) \\ \boxed{p} &= \frac{1}{n} \sum_{i=1}^n x_i \end{aligned}$$

2.

$$\mathbb{E}[\hat{p}] = \mathbb{E}\left[\frac{1}{n} \sum_{i=1}^n x_i\right] = \frac{1}{n} (\mathbb{E}[x_i]) = \frac{1}{n} np = p$$

Therefore, \hat{p} is an unbiased estimator of p

3. Expected square error of \hat{p} in terms of p can be written as

$$\begin{aligned} \mathbb{E}[(p - L(p))^2] &= \mathbb{E}[(p - \hat{p})^2] = \mathbb{E}[p^2 - 2p\hat{p} + \hat{p}^2] = \mathbb{E}[p^2] - \mathbb{E}[2p\hat{p}] + \mathbb{E}[\hat{p}^2] \\ &= p^2 - \frac{2p}{n} \sum_{i=1}^n \mathbb{E}[x_{i=1}^n] + \frac{1}{n} \sum_{i=1}^n \mathbb{E}[x_{i=1}^n] \\ &= p^2 - \frac{2np^2}{n} + \frac{np}{n^2}((n-1)p + 1) \\ &= p^2 - 2p^2 - \frac{p^2}{n} + \frac{p}{n} + p^2 \\ \boxed{\text{MSE}(\hat{p})} &= \frac{p(1-p)}{n} \end{aligned}$$

4. We know that $p \in \{\frac{1}{4}, \frac{3}{4}\}$ and $n = 3$. $R(\theta, \delta) = \frac{p(1-p)}{n}$. For $p \in \{\frac{1}{4}, \frac{1}{2}, \frac{3}{4}\}$, we have the following values for $R(\theta, \delta) = \{0.0625, 0.0833, 0.0625\}$. \hat{p} would be an inadmissible estimator of p if there exists an $R(\theta, \delta)$ that gives a better estimate of MSE . Intuitively, the best estimator of p would have $MSE = 0$ (lowest uncertainty). Hence, in this case, \hat{p} would be an inadmissible estimator for p .

5 Problem 5

$$\mu_{\text{MLE}} = \frac{1}{n} \sum_{i=1}^n x_i$$

$$\sigma_{\text{MLE}}^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \mu_{\text{MLE}})^2$$

•

$$\begin{aligned} \mathbb{E}[\sigma_{\text{MLE}}^2] &= \mathbb{E} \left[\frac{1}{n} \sum_{i=1}^n (x_i - \mu_{\text{MLE}})^2 \right] \\ &= \mathbb{E} \left[\frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2 \right] = \mathbb{E} \left[\frac{1}{n} \sum_{i=1}^n (x_i^2 - 2x_i\bar{x} + \bar{x}^2) \right] = \frac{1}{n} \mathbb{E} \left[\sum_{i=1}^n (x_i^2) - 2n\bar{x}^2 + n\bar{x}^2 \right] \\ &= \frac{1}{n} \mathbb{E} \left[\sum_{i=1}^n (x_i^2) - n\bar{x}^2 \right] = \frac{1}{n} \mathbb{E} \left[\sum_{i=1}^n x_i^2 \right] - \mathbb{E}[\bar{x}^2] = \mathbb{E}[x_i^2] - \mathbb{E}[\bar{x}^2] \end{aligned}$$

We know that $\text{Var}[X] = \mathbb{E}[x^2] - \mathbb{E}[x]^2$ from the definition of variance.

$$\begin{aligned} \mathbb{E}[x_i^2] - \mathbb{E}[\bar{x}^2] &= \sigma_x^2 + \mathbb{E}[x_i]^2 - \sigma_x^2 - \mathbb{E}[x_i]^2 = \sigma_x^2 - \sigma_{\bar{x}}^2 \\ &= \sigma_x^2 - \text{Var}[\bar{x}] = \sigma_x^2 - \text{Var} \left[\frac{1}{n} \sum_{i=1}^n x_i \right] \\ &= \sigma_x^2 - \left(\frac{1}{n} \right)^2 \text{Var} \left[\sum_{i=1}^n x_i \right] \end{aligned}$$

Now,

$$\sigma_x^2 - \left(\frac{1}{n} \right)^2 \text{Var} \left[\sum_{i=1}^n x_i \right] = \sigma_x^2 - \left(\frac{1}{n} \right)^2 n\sigma_x^2 = \frac{n-1}{n} \sigma_x^2 < \sigma_x^2$$

Thus, σ_{MLE}^2 is an unbiased estimator of σ^2 .

•

$$\begin{aligned} \mathbb{E}[\hat{\sigma}^2] &= \mathbb{E} \left[\frac{1}{n} \sum_{i=1}^n (x_i - \mu)^2 \right] = \frac{1}{n} \sum_{i=1}^n \mathbb{E}[(x_i - \mu)^2] = \frac{1}{n} \sum_{i=1}^n \mathbb{E}[x_i^2 - 2x_i\mu + \mu^2] \\ &= \frac{1}{n} \left[\sum_{i=1}^n \mathbb{E}[x_i^2] - 2 \sum_{i=1}^n \mathbb{E}[x_i\mu] + \sum_{i=1}^n \mathbb{E}[\mu^2] \right] = \frac{1}{n} [n\sigma^2 + n^2\mu^2 + n^2\mu^2 - 2n^2\mu^2] = \frac{1}{n} [n\sigma^2] = \sigma^2 \end{aligned}$$

Hence, $\hat{\sigma}^2$ is an unbiased estimator of σ^2 .

6 Problem 6

Consider a covariance matrix Σ for any random variable X . A covariance matrix is symmetric by definition. Next, consider a vector $Z \rightarrow (z_1, \dots, z_n)$ in n space.

$$\begin{aligned} Z^T &= Z^T \mathbb{E}[(x - \mathbb{E}[x])(x - \mathbb{E}[x])^T] \\ &= \mathbb{E}[(x - \mathbb{E}[x])^T \cdot Z)^T ((x - \mathbb{E}[x])^T \cdot Z)] \\ &= \mathbb{E}[\underbrace{((x - \mathbb{E}[x])^T \cdot Z)^T}_{\text{Non-negative}}] \geq 0 \end{aligned}$$

Since, the value of the term within the expectation is non-negative, the expectation must be non-negative. Therefore, $Z^T \Sigma Z \geq 0$. Now, if X were to be a Gaussian random variable, we know that the covariance matrix Σ must be invertible and full rank. Also, $Z^T \Sigma Z > 0$ and is positive definite.

7 Problem 7

$$\|X^T w - y\|^2 = \underbrace{(w - \hat{w})^T X^T X (w - \hat{w})}_{(x - M^{-1}b)^T x^T x (x - M^{-1}b)} + \|X^T \hat{w} - y\|^2$$

We can show that $\|X^T w - y\|^2 - \|X^T \hat{w} - y\|^2 = (w - \hat{w})^T X^T X (w - \hat{w})$

Let $x = w, M = X^T X, b = X^T y$. Then $(w - \hat{w})^T X^T X (w - \hat{w})$ can be written as $(x - M^{-1}b)^T x^T x (x - M^{-1}b)$. Also, we know that

$$(x - M^{-1}b)^T x^T x (x - M^{-1}b) = x^T M x - 2b^T x + b^T M^{-1}b$$

Therefore

$$\begin{aligned} (w - \hat{w})^T X^T X (w - \hat{w}) &= w^T X^T X w - 2(X^T y)^T w + (X^T y)^T (X^T X)^{-1} (X^T y) \\ &= w^T X^T X w - 2(X^T y)^T w + (X^T y)^T \hat{w} (X^T y)^{-1} (X^T y) \end{aligned}$$

Next,

$$\begin{aligned} \|X^T w - y\|^2 - \|X^T \hat{w} - y\|^2 &= (X^T w)^T (X^T w) - 2y^T (X^T w) + y^T y - (X^T \hat{w})^T (X^T \hat{w}) + 2y^T (X^T \hat{w}) - y^T y \\ &= (X^T w)^T (X^T w) - (X^T \hat{w})^T (X^T \hat{w}) - 2y^T (X^T w) + 2y^T (X^T \hat{w}) \end{aligned}$$